

# Class 8 Mini Project

Emily Chase (PID A14656894)

## Table of contents

Background . . . . .	1
Data Import . . . . .	1
PCA . . . . .	3
Calculate variance of each component . . . . .	11
Communicating PCA results . . . . .	12
Hierarchical clustering . . . . .	13
Using different methods . . . . .	14
Combining PCA and Clustering . . . . .	16

## Background

Extend what you've learned by combining PCA as a preprocessing step to clustering using data that consist of measurements of cell nuclei of human breast masses.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets".

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

## Data Import

```
fna.data <- "WisconsinCancer.csv"

# We'll use the identifiers as names (ie don't leave them in the df) so that it doesn't accidentally
wisc.df <- read.csv(fna.data, row.names=1)
colnames(wisc.df)
```

```

[1] "diagnosis"                  "radius_mean"
[3] "texture_mean"               "perimeter_mean"
[5] "area_mean"                  "smoothness_mean"
[7] "compactness_mean"            "concavity_mean"
[9] "concave.points_mean"        "symmetry_mean"
[11] "fractal_dimension_mean"     "radius_se"
[13] "texture_se"                 "perimeter_se"
[15] "area_se"                    "smoothness_se"
[17] "compactness_se"              "concavity_se"
[19] "concave.points_se"          "symmetry_se"
[21] "fractal_dimension_se"       "radius_worst"
[23] "texture_worst"              "perimeter_worst"
[25] "area_worst"                 "smoothness_worst"
[27] "compactness_worst"           "concavity_worst"
[29] "concave.points_worst"        "symmetry_worst"
[31] "fractal_dimension_worst"

```

Since our analysis is going to be unsupervised, we are going to take out the “answers” (the professional diagnosis)

```

# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
# save the excluded column
diagnosis <- as.factor(wisc.df$diagnosis)

head(wisc.data[,1:3])

```

	radius_mean	texture_mean	perimeter_mean
842302	17.99	10.38	122.80
842517	20.57	17.77	132.90
84300903	19.69	21.25	130.00
84348301	11.42	20.38	77.58
84358402	20.29	14.34	135.10
843786	12.45	15.70	82.57

Q1. How many observations are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

There are 569 observations/samples/patients in the data set.

Q2. How many of the observations have a malignant diagnosis?

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

```
table(wisc.df$diagnosis)
```

B	M
357	212

There are `sum(wisc.df$diagnosis == "M")` malignant diagnoses.

Q3. How many variables/features in the data are suffixed with `_mean`?

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

There are `length(grep("_mean", colnames(wisc.df)))` columns that give a mean summary value.

## PCA

The main function in base R for PCA is called `prcomp()`. It has arguments `x`, `scale=F`, `center=T`. `Center` is about whether they should zero center the data (subtract by the mean to have mean 0). `scale` is about whether the data should be scaled to have unit variance. PCA will unfairly be influenced by the columns that have the most variance. Some columns will have more variance just because you picked a different unit.

**You almost always want to scale your data.** This means you'll always set the argument `scale=TRUE`.

It is important to check if the data need to be scaled before performing PCA. Recall two common reasons for scaling data include:

- The input variables use different units of measurement.
- The input variables have significantly different variances.

Check the mean and standard deviation of the features (i.e. columns) of the wisc.data to determine if the data should be scaled. Use the colMeans() and apply() functions like you've done before.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03

	symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03		2.646071e-03	4.833242e+00
	texture_worst	perimeter_worst	area_worst
6.146258e+00		3.360254e+01	5.693570e+02
	smoothness_worst	compactness_worst	concavity_worst
2.283243e-02		1.573365e-01	2.086243e-01
	concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02		6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Let's make our main result figure - the "PC Plot" or "score plot", "ordination plot", etc


```
head(wisc.pr$x)
```

	PC1	PC2	PC3	PC4	PC5	PC6
842302	-9.184755	-1.946870	-1.1221788	3.6305364	1.1940595	1.41018364
842517	-2.385703	3.764859	-0.5288274	1.1172808	-0.6212284	0.02863116

84300903	-5.728855	1.074229	-0.5512625	0.9112808	0.1769302	0.54097615
84348301	-7.116691	-10.266556	-3.2299475	0.1524129	2.9582754	3.05073750
84358402	-3.931842	1.946359	1.3885450	2.9380542	-0.5462667	-1.22541641
843786	-2.378155	-3.946456	-2.9322967	0.9402096	1.0551135	-0.45064213
	PC7	PC8	PC9	PC10	PC11	PC12
842302	2.15747152	0.39805698	-0.15698023	-0.8766305	-0.2627243	-0.8582593
842517	0.01334635	-0.24077660	-0.71127897	1.1060218	-0.8124048	0.1577838
84300903	-0.66757908	-0.09728813	0.02404449	0.4538760	0.6050715	0.1242777
84348301	1.42865363	-1.05863376	-1.40420412	-1.1159933	1.1505012	1.0104267
84358402	-0.93538950	-0.63581661	-0.26357355	0.3773724	-0.6507870	-0.1104183
843786	0.49001396	0.16529843	-0.13335576	-0.5299649	-0.1096698	0.0813699
	PC13	PC14	PC15	PC16	PC17	
842302	0.10329677	-0.690196797	0.601264078	0.74446075	-0.26523740	
842517	-0.94269981	-0.652900844	-0.008966977	-0.64823831	-0.01719707	
84300903	-0.41026561	0.016665095	-0.482994760	0.32482472	0.19075064	
84348301	-0.93245070	-0.486988399	0.168699395	0.05132509	0.48220960	
84358402	0.38760691	-0.538706543	-0.310046684	-0.15247165	0.13302526	
843786	-0.02625135	0.003133944	-0.178447576	-0.01270566	0.19671335	
	PC18	PC19	PC20	PC21	PC22	
842302	-0.54907956	0.1336499	0.34526111	0.096430045	-0.06878939	
842517	0.31801756	-0.2473470	-0.11403274	-0.077259494	0.09449530	
84300903	-0.08789759	-0.3922812	-0.20435242	0.310793246	0.06025601	
84348301	-0.03584323	-0.0267241	-0.46432511	0.433811661	0.20308706	
84358402	-0.01869779	0.4610302	0.06543782	-0.116442469	0.01763433	
843786	-0.29727706	-0.1297265	-0.07117453	-0.002400178	0.10108043	
	PC23	PC24	PC25	PC26	PC27	
842302	0.08444429	0.175102213	0.150887294	-0.201326305	-0.25236294	
842517	-0.21752666	-0.011280193	0.170360355	-0.041092627	0.18111081	
84300903	-0.07422581	-0.102671419	-0.171007656	0.004731249	0.04952586	
84348301	-0.12399554	-0.153294780	-0.077427574	-0.274982822	0.18330078	
84358402	0.13933105	0.005327110	-0.003059371	0.039219780	0.03213957	
843786	0.03344819	-0.002837749	-0.122282765	-0.030272333	-0.08438081	
	PC28	PC29	PC30			
842302	-0.0338846387	0.045607590	0.0471277407			
842517	0.0325955021	-0.005682424	0.0018662342			
84300903	0.0469844833	0.003143131	-0.0007498749			
84348301	0.0424469831	-0.069233868	0.0199198881			
84358402	-0.0347556386	0.005033481	-0.0211951203			
843786	0.0007296587	-0.019703996	-0.0034564331			

```
library(ggplot2)
```

```
ggplot(wisc.pr$x) +aes(x=PC1, y=PC2, col=diagnosis) + geom_point()
```



```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					

```

Standard deviation      0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion  1.00000 1.00000

```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% of the original variance is described by PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


According to cumulative proportion of the variance, we see that the cumulative variance crosses the 70% threshold at PC3, so 3 PCs are required.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?


7 PCs gets us to 91% cumulative variance.

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

```
ggplot(wisc.pr$x) + aes(x=PC1, y=PC2) + geom_text(aes(label = rownames(wisc.pr$x)))
```




```
ggplot(wisc.pr$x) + aes(x=PC1, y=PC2) + geom_point()
```



Using text makes the overlapping data impossible to parse, but using points shows the spread a little more clearly. Still can't see the groupings well yet though because it's all in one color.


Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) + aes(x=PC1, y=PC3) + geom_point()
```



The data looks even messier using PC3 instead of PC2, because PC2 explains more variance than PC3.

```
ggplot(wisc.pr$x) + aes(x=PC1, y=PC3, col=diagnosis) + geom_point()
```



### Calculate variance of each component


```
pr.var <- wisc.pr$sdev^2
head(pr.var)

[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357


# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)
pve

[1] 4.427203e-01 1.897118e-01 9.393163e-02 6.602135e-02 5.495768e-02
[6] 4.024522e-02 2.250734e-02 1.588724e-02 1.389649e-02 1.168978e-02
[11] 9.797190e-03 8.705379e-03 8.045250e-03 5.233657e-03 3.137832e-03
[16] 2.662093e-03 1.979968e-03 1.753959e-03 1.649253e-03 1.038647e-03
[21] 9.990965e-04 9.146468e-04 8.113613e-04 6.018336e-04 5.160424e-04
[26] 2.725880e-04 2.300155e-04 5.297793e-05 2.496010e-05 4.434827e-06

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
barplot(pve)
```



## Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

```
head(wisc.pr$rotation[,1:3], 10)
```

	PC1	PC2	PC3
radius_mean	-0.21890244	0.23385713	-0.008531243
texture_mean	-0.10372458	0.05970609	0.064549903
perimeter_mean	-0.22753729	0.21518136	-0.009314220
area_mean	-0.22099499	0.23107671	0.028699526
smoothness_mean	-0.14258969	-0.18611302	-0.104291904
compactness_mean	-0.23928535	-0.15189161	-0.074091571
concavity_mean	-0.25840048	-0.06016536	0.002733838
concave.points_mean	-0.26085376	0.03476750	-0.025563541
symmetry_mean	-0.13816696	-0.19034877	-0.040239936
fractal_dimension_mean	-0.06436335	-0.36657547	-0.022574090


The loading for concave.points\_mean is -0.26085376.

## Hierarchical clustering

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist, method="complete")

plot(wisc.hclust)
abline(h=18, col="red", lty=2)
```



```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165

```

2   2   5
3 343  40
4   0   2

```

### Using different methods


Q12. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning.

```

data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)

wisc.hclust1 <- hclust(data.dist, method="complete")
plot(wisc.hclust1)

```




**data.dist**  
**hclust (\*, "complete")**

```

wisc.hclust2 <- hclust(data.dist, method="single") #weird and curvy
plot(wisc.hclust2)

```


## Cluster Dendrogram



```
data.dist  
hclust (*, "single")
```


```
wisc.hclust3 <- hclust(data.dist, method="average") # shows three distinctish groups  
plot(wisc.hclust3)
```

## Cluster Dendrogram



```
data.dist  
hclust (*, "average")
```

```
wisc.hclust4 <- hclust(data.dist, method="ward.D2") # this is the clearest
plot(wisc.hclust4)
```




I like `ward.D2` the best. It shows the clearest differences between groups.

### Combining PCA and Clustering

Use knowledge from our PCA (eg our scree plots!) and using it to inform our dendrogram

```
d <- dist(wisc.pr$x[,1:3]) # we know from our scree plot that we should just choose the first
wisc.pr.hclust <- hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
abline(h=70, col="red")
```

## Cluster Dendrogram



```
d  
hclust (*, "ward.D2")
```

Get my cluster membership

```
grps <- cutree(wisc.pr.hclust, h=70)  
table(grps)
```

```
grps  
1 2  
203 366
```

```
table(diagnosis)
```

```
diagnosis  
B M  
357 212
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

Is my clustering capturing this B vs M? Let's compare our clusters (`grps`) in `diagnosis`. Make a wee “cross-table”.

```
table(grps,diagnosis)
```

	diagnosis	
grps	B	M
1	24	179
2	333	33

True Positive = 179 False Positive = 24

True Negative = 333 False Negative = 33

Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

This comparison is a little trickier because it has four categories vs the binary diagnosis. But based on where the majority of the values lie, there are fewer false positives but more false negatives. I personally think false negatives are even scarier than false positives (false positives can mean further testing but false negatives can mean that no further investigation will happen)